



Figure 1: An overview of the TAPAS approach. A dense table retriever scores the question against all tables and outputs the top  $K$  tables ( $K = 1$  in this example), and a reader selects the answer from the top table.